# CLARITY - SemEval 2026 Challenge


**LLMSE 2026 Course Project - Politecnico di Torino**

**Authors**

Stefano Zizzi 346595
Alessio Perrotti 346737
Riccardo Vaccari 348856
Davide Candela 347245
Luca Lodesani 346978



## Abstract

Detecting political evasion is a complex natural language processing task that requires understanding subtle rhetorical strategies rather than just lexical matching. The SemEval-2026 CLARITY challenge formalizes this problem by requiring the classification of political responses into distinct clarity and evasion categories. However, severe class imbalance and the inherently subjective nature of political discourse make classification difficult, particularly in resource-constrained computational environments. In this paper, we present an empirical, resource-efficient approach to political evasion detection by comparing lightweight bidirectional encoders with quantized generative Large Language Models (LLMs). Our results demonstrate that injecting structural taxonomy into a smaller, specialized architecture is a highly effective and computationally accessible alternative to relying on the zero-shot reasoning of larger models.


## 1 Introduction

In political communication, direct answers are often the exception rather than the rule. Equivocation, dodging, and deflection are strategic tools used to navigate complex interviews without committing to highly scrutinized positions. Historically, measuring this phenomenon at scale has been difficult due to the subjective nature of human intent. The SemEval-2026 CLARITY challenge addresses this by providing a framework to objectively classify response clarity and rhetorical evasion in political interviews. The challenge requires systems to evaluate question-answer pairs at two levels of granularity: a macro-level assessment of overall clarity (Task 1) and a fine-grained categorization of specific evasive strategies (Task 2).

Developing robust systems for this task presents two significant hurdles. First, real-world political discourse inherently suffers from severe class imbalance; ambiguous responses heavily outnumber explicit refusals or clear replies. Second, relying on massive, parameter-heavy architectures to capture these nuanced rhetorical shifts is often computationally prohibitive, limiting the accessibility of computational discourse analysis. This paper embraces these constraints, focusing specifically on resource-efficient methodologies for evasion detection. To address these challenges, we present a comprehensive empirical evaluation of computationally light architectures, contrasting discriminative and generative paradigms. Our investigation is divided into two primary tracks. First, we explore the limits of lightweight bidirectional encoders (DistilBERT and RoBERTa). To prevent model collapse under severe class imbalance, we introduce an "Enhanced" pipeline that combines a novel Smart Resampling technique (KMeans downsampling paired with T5 paraphrase upsampling) with the explicit injection of zero-shot rhetorical tone metadata prior to tokenization. Second, we investigate the efficacy of quantized generative Large Language Models (LLMs) through an ablation study on model scale versus adaptation strategy. We specifically contrast the performance of a larger reasoning engine against a smaller, highly specialized one. To test heuristic reasoning, we apply complex prompt engineering—including Chain-of-Thought (CoT) and few-shot learning to a quantized Llama-3.1-8B-Instruct model. We benchmark this against

Figure 1: Macro F1 progression across few-shot strategies.

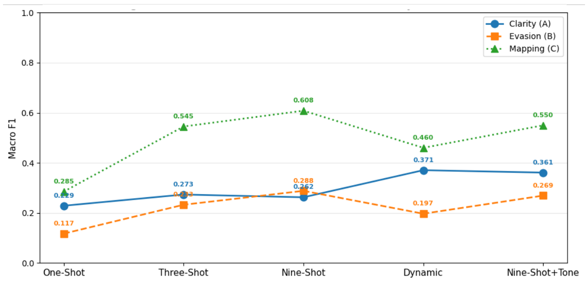

parameter-efficient fine-tuning (QLoRA), where we inject the task's structural taxonomy directly into the weights of a more compact Llama-3.2-3B-Instruct model. Our experimental results demonstrate that predicting fine-grained evasion strategies and mapping them deterministically to macro-clarity labels provides a superior supervision signal. Notably, our findings show that structural adaptation on the smaller 3B model yields the most robust performance, outperforming both the robustly balanced encoders and the larger, prompt-engineered 8B architecture.

## 2 Background

### 2.1 Equivocation and Political Speech Analysis

In social and political sciences, equivocation is defined as a strategic communicative act used to avoid answering questions directly while maintaining rhetorical control (Dillon, 1990). Foundational frameworks (Harris, 1991; Bull, 2003) established comprehensive typologies to categorize responses into direct replies, intermediate answers, and various non-reply evasion techniques (e.g., dodging, attacking the question).

In Natural Language Processing (NLP), this phenomenon has historically been approached through the lenses of question answerability (Rajpurkar et al., 2018) or deceptive intent detection (Girlea, 2017). However, intent-based frameworks often suffer from high subjectivity. To address this, Thomas et al. (2024) introduced the QEvasion dataset, shifting the focus to a deterministic, two-level taxonomy that evaluates objective response clarity rather than perceived speaker intent, providing a robust foundation for supervised learning.

### 2.2 Class Imbalance and Data Augmentation

A persistent challenge in modeling political discourse is the severe class imbalance inherent to real-world interviews, where ambiguous responses heavily outnumber explicit non-replies. Traditional text augmentation techniques, such as synonym replacement or random word deletion (Wei and Zou, 2019), often fail to preserve the delicate rhetorical structures required for evasion detection. To mitigate this without introducing synthetic noise, recent methodologies favor model-based resampling. Techniques such as semantic clustering via sentence embeddings (Reimers and Gurevych, 2019) allow for the diverse downsampling of majority classes, while Seq2Seq models like T5 (Raffel et al., 2023) enable the generation of high-quality, linguistically varied paraphrases for minority classes, reducing the risk of model overfitting.

### 2.3 Transformer Encoders and PEFT Decoders

The task of response clarity classification relies heavily on capturing contextual relationships between questions and answers. Transformer-based encoders, particularly RoBERTa (Liu et al., 2019), have demonstrated strong performance in sequence classification tasks due to their bidirectional attention mechanisms.

Conversely, the advent of Large Language Models (LLMs) like Llama 3 (Grattafiori et al., 2024) has introduced advanced reasoning capabilities. Since full-parameter training of such models is computationally expensive, Parameter-Efficient Fine-Tuning (PEFT) methods—specifically Low-Rank Adaptation (LoRA) (Hu et al., 2022)— have become the standard for domain adaptation. By updating only a small subset of weights, LoRA enables generative models to learn complex, fine-grained taxonomic classifications with high efficiency.

## 3 System overview

This section details the architectural choices, data curation strategies, and modeling paradigms implemented to detect political evasion. Given the complexity of assessing response clarity and rhetorical strategies, our system relies on a multi-stage pipeline. The first segment (Sections 3.1–3.4) covers dataset formulation, semantic resampling, feature engineering, and the training of discriminative transformer encoders. The second segment (Sections 3.5 and 3.6) explores the application of generative Large Language Models (LLMs) via parameter-efficient fine-tuning (PEFT) and prompt chaining frameworks.

### 3.1 Dataset Description and Task Formulation

The experimental pipeline is built upon the QEvasion dataset (AilsntuaLab, 2024), which formulates political evasion detection as a supervised classification problem over question–answer (QA) pairs. The objective is to identify whether a response appropriately addresses a given question and, when evasion occurs, to characterize the underlying rhetorical strategy. The dataset therefore

Table 1: Class distribution for Task 2 in training set

| Evasion Strategy | Train | % |
|---|---|---|
| Explicit | 1051 | 30.5% |
| Dodging | 704 | 20.5% |
| Implicit | 488 | 14.2% |
| General | 386 | 11.2% |
| Deflection | 381 | 11.0% |
| Declining to answer | 145 | 4.2% |
| Claims ignorance | 119 | 3.5% |
| Clarification | 92 | 2.7% |
| Partial half-answer | 79 | 2.3% |

defines two complementary tasks operating at different levels of granularity:

- **Task 1 (Clarity Classification)**: A coarse-grained three-class classification task aimed at identifying the overall communicative outcome of the response. The possible labels are *Clear Reply*, *Clear Non-Reply*, and *Ambivalent*, capturing whether the answer provides a direct response, avoids the question, or exhibits mixed characteristics.

- **Task 2 (Evasion Strategy Classification)**: A fine-grained nine-class classification task focused on identifying the specific rhetorical mechanism employed when an evasive behavior is detected. The target categories include strategies such as *Dodging*, *Deflection*, and *Claims ignorance*, among others (see Table 1).

While Task 1 provides a high-level characterization of response clarity, Task 2 requires a deeper understanding of pragmatic and rhetorical patterns. In particular, the latter presents a significant challenge due to both the severe class imbalance (see Table 1) and the higher degree of annotator subjectivity involved in distinguishing between closely related evasion strategies. The most frequent class is over ten times larger than the rarest categories, making effective representation learning and robust classification particularly challenging.

### 3.2 Data Balancing: Smart Resampling

To prevent the models from collapsing into majority-class prediction (Mode Collapse) due to the severe imbalance in the evasion strategies, we implemented a "Smart Resampling" pipeline combining semantic downsampling and paraphrase upsampling.

For downsampling the over-represented classes (e.g., Ambivalent), we employed a KMeans clustering approach. Given a target class size *T*, we extract sentence embeddings using the `all-MiniLM-L6-v2` model and compute *T* cluster centroids. We retain only the real samples closest to these centroids, ensuring optimal semantic diversity while discarding redundant political rhetoric.

For upsampling minority classes (e.g., Clear Non-Reply), we utilized a T5-based sequence-to-sequence model (`humarin/chatgpt_paraphraser_on_T5_base`) to generate synthetic paraphrases. To prevent overfitting from robotic artifacts, generation was strictly capped to a 1:1 ratio (maximum one paraphrase per original example).

Furthermore, we implemented a semantic filter computing the cosine similarity between the original sentence $S_{orig}$ and the paraphrase $S_{new}$:

$$\text{Sim}(S_{orig}, S_{new}) = \frac{E(S_{orig}) \cdot E(S_{new})}{\|E(S_{orig})\| \, \|E(S_{new})\|} \tag{1}$$

Only paraphrases falling within a similarity threshold of $[0.60, 0.90]$ were accepted, guaranteeing syntactic novelty while strictly preserving the original evasion label.

### 3.3 Feature Engineering: Rhetorical Tone

A key methodological novelty in our system is the extraction and injection of rhetorical metadata into the text prior to tokenization. Recognizing that evasion is heavily dependent on the speaker's attitude rather than purely lexical cues, we engineered a specific feature to capture communicative intent.

We utilized a Zero-Shot classification pipeline based on `facebook/bart-large-mnli` to evaluate each answer $A_i$. The classifier computes a probability distribution over three mutually exclusive candidate labels $L$ (Assertive, Guarded, Dismissive), which were specifically chosen to represent the spectrum of political rhetoric. The assigned tone $T_i$ is determined by selecting the label with the highest predicted probability:

$$T_i = \arg\max_{l \in L} P(l \mid A_i, \text{BART}) \tag{2}$$

To leverage this metadata, we explicitly prepended the extracted tone to the input sequence. For the transformer encoders, the final input provided to the tokenizer was structured as follows:

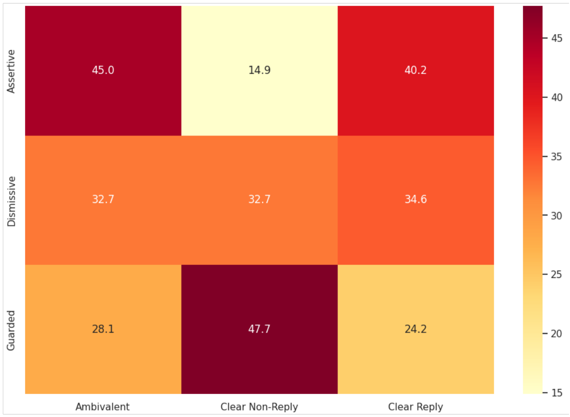

Figure 2: Correlation between rhetorical tone and clarity classes.

```
[CLS] Question [SEP] Tone: Extracted Tone |
        Answer: Answer text [SEP]
```

This formatting forces the bidirectional attention mechanism to condition the question-answer representation on the injected rhetorical hint, providing a semantic anchor during classification.

**Empirical Validation.** To justify the selection of these specific tone labels, we analyzed their correlation with the ground-truth clarity classes on the balanced training set. As illustrated in Figure 2, the predicted tone serves as a highly discriminative prior:

- **Guarded**: Correlates strongly (47.7%) with explicit refusals (Clear Non-Reply), effectively isolating evasive caution.
- **Dismissive**: Exhibits a balanced distribution, capturing blunt responses that act either as short refusals (Clear Non-Reply) or brief factual confirmations (Clear Reply).
- **Assertive**: Although expected to indicate clarity, an assertive tone predominantly maps to Ambivalent responses (45.0%). This suggests that the model successfully captures the political strategy of delivering evasive or off-topic answers with a confident and factual demeanor.

### 3.4 Model Architecture and Training Procedure

For the sequence classification task, we benchmarked two distinct bidirectional transformer encoders, evaluating the trade-off between predictive performance and computational efficiency:

- **DistilBERT-base-uncased**: A parameter-efficient architecture containing 6 transformer layers and 66 million parameters. Through knowledge distillation, it retains approximately 97% of the original BERT capabilities while reducing the parameter count by 40%. It serves as our lightweight baseline, providing faster training and inference times suitable for resource-constrained environments.
- **RoBERTa-base**: A robustly optimized architecture comprising 12 transformer layers and 125 million parameters. Unlike the original BERT training procedure, RoBERTa employs dynamic masking during pre-training on a significantly larger corpus (160GB of text) and removes the Next Sentence Prediction (NSP) objective. These modifications allow the model to capture deeper semantic relationships and subtle rhetorical nuances, making it our primary encoder choice for maximizing performance on complex evasion detection tasks.

#### 3.4.1 Training Procedure and Optimization

To ensure a fair and stable comparison across architectures, we adopted a standardized fine-tuning protocol. All encoders were optimized using the AdamW optimizer with a weight decay coefficient of 0.01 to mitigate excessive parameter updates. Models were trained for a maximum of 5 epochs.

Given the effectiveness of the proposed *Smart Resampling* pipeline (Section 3.2), the resulting training distribution achieved sufficient class balance. Therefore, more advanced loss formulations, such as focal loss, were not required, and standard Cross-Entropy Loss was adopted:

$$\mathcal{L}_{CE} = -\sum_{i=1}^{C} y_i \log(\hat{y}_i) \tag{3}$$

where $C$ denotes the number of classes, $y_i$ is the ground-truth label, and $\hat{y}_i$ represents the predicted probability for class $i$.

Furthermore, considering the sensitivity of transformer attention layers during the initial optimization steps, we employed a linear learning rate scheduler with a warmup ratio of 0.1. This strategy allowed the randomly initialized classification heads to progressively stabilize, enabling both RoBERTa and DistilBERT to effectively learn from the injected rhetorical metadata without training instability.

### 3.5 Generative LLM Prompt Engineering

To establish a strong heuristic reasoning baseline prior to parameter-efficient fine-tuning, we explored training-free prompt engineering techniques using a quantized generative model (`Llama-3.1-8B-Instruct`). Running in 4-bit precision (NF4) to fit resource constraints, this configuration allowed us to test the model's analytical capabilities out-of-the-box. Our investigation followed a chronological progression from Zero-Shot and Chain-of-Thought (CoT) reasoning to more complex Few-Shot In-Context Learning strategies.

**Chain-of-Thought.** We formulated the evasion detection task as a prompt-based generation problem, incrementally enriching the context provided to the LLM:

- **Zero-Shot:** We established a baseline using the exact classification prompt provided in the appendix of the original CLARITY paper (Thomas et al., 2024), which strictly provides the class definitions without any explicit reasoning instructions.

- **Chain-of-Thought (CoT):** To guide the model's analytical process, we formulated a novel, structured 3-step reasoning chain. Before outputting a label, the LLM was forced to: (1) summarize the core information requested by the journalist, (2) concretely describe the politician's actual rhetorical maneuver, and (3) justify which label fits best.

- **Taxonomic Table Injection (CoT + Table):** The original paper (Thomas et al., 2024) relies on a detailed taxonomy table containing definitions and examples, used exclusively as a guideline for human annotators. We injected this human-annotation table directly into the CoT prompt to explicitly define class boundaries for the LLM.

- **Retrieval and Heuristics (CoT + Table + RAG):** We augmented the context in two ways. First, to prevent internal model bias regarding specific political entities or locations, we queried Wikipedia for capitalized named entities present in the question, appending their summaries to the prompt. Second, we implemented a heuristic question-type classifier (e.g., mapping "Why" to a medium-length explanation, or "When" to a short temporal reference). This injected an "expected length" and a complexity note into the prompt, preventing the model from erroneously penalizing genuine but lengthy answers or misclassifying short direct replies.

- **Rhetorical Tone (CoT + Table + Tone):** Recognizing that evasion is heavily characterized by attitude, we integrated a pre-calculated rhetorical tone (Assertive, Guarded, Dismissive) into the prompt as a critical structural hint. We explicitly instructed the LLM to use this tone during an additional CoT reasoning step, conditioning its final decision on the communicative intent rather than just the semantic content.

**Few-Shot.** Building upon the most successful CoT frameworks, we evaluated the impact of In-Context Learning by providing the model with solved examples from the training set. To avoid introducing length bias in the model's generations, we conducted an analysis of the answer length distribution in the training set. We strategically selected "representative" examples whose character length was closest to the median length of their respective class, ensuring that the LLM was not skewed by extreme outliers.

- **Static Few-Shot:** We progressively increased the context size, testing One-Shot (a single representative example of a clear reply), Three-Shot (one median-length example per macro-clarity class), and Nine-Shot (one median-length example per fine-grained evasion category).

- **Dynamic Few-Shot (Semantic Retrieval):** To maximize context relevance, we hypothesized that dynamically retrieving the most similar training examples would yield better results. We built a FAISS vector index using the `all-MiniLM-L6-v2` embedding model. For each test instance, we dynamically retrieved the nearest $k = 3$ neighbors based on cosine similarity, explicitly enforcing a diversity constraint (maximum one example per label) to ensure a balanced in-context learning environment.

- **Nine-Shot + Tone:** Finally, we combined the most extensive static context (Nine-Shot) with the previously validated rhetorical tone injection methodology, providing the model with both behavioral examples and explicit intent conditioning.

### 3.6 Generative LLM Fine-Tuning for Evasion Classification

In addition to the other classifiers, we explored an instruction-tuned generative approach based on `meta-llama/Llama-3.2-3B-Instruct`.

The prompt used during fine-tuning frames the problem as an expert political communication analysis. For each example, the model receives the journalist's question and the politician's answer, and is required to return a valid JSON object containing exactly one evasion category. This structured-output format was chosen to reduce ambiguity at inference time and to make the predictions directly post-processable.

The model was loaded with 4-bit quantization using NF4 quantization and double quantization. Fine-tuning was performed with QLoRA/-DoRA adapters on the attention projection modules `q_proj`, `k_proj`, `v_proj`, and `o_proj`. We used rank $r = 16$, LoRA alpha 32, dropout 0.05, an AdamW 8-bit optimizer, a learning rate of $2 \cdot 10^{-4}$, cosine scheduling, maximum sequence length 512, and one training epoch. This setup keeps the number of trainable parameters small while still allowing the model to specialize on the rhetorical taxonomy of the task.

## 4 Experimental Results

We evaluate our proposed models on three experimental settings: directly predicting the 3-class clarity labels (Task 1), directly predicting the 9 fine-grained evasion strategies (Task 2), and a hierarchical approach that first predicts evasion strategies and then maps them deterministically to clarity labels.

We compare two configurations across encoder-based models. The **Baseline** applies standard fine-tuning of DistilBERT-base-uncased or RoBERTa-base on the original, unmodified training split—no data augmentation, no feature engineering, and standard cross-entropy loss over the imbalanced class distribution. The **Enhanced** configuration extends this setup with two complementary modules: *Smart Resampling* (Section 3.2) and *Tone Analysis* (Section 3.3).

All evaluations use the same held-out test set of 308 question–answer pairs; Macro-F1 is the primary metric, as the test set is class-imbalanced.

### 4.1 Task 1 (Direct Clarity)

Models predict the three clarity classes (*Clear Reply*, *Clear Non-Reply*, *Ambivalent*) directly. Results are reported in Table 2.

Table 2: Task 1 (Direct Clarity) results for all systems.

| Model Family | Configuration | Acc. | Macro-F1 | Weighted-F1 |
|---|---|---|---|---|
| Encoders | Baseline (DistilBERT) | 0.643 | 0.545 | 0.636 |
| | Enhanced (DistilBERT) | 0.601 | 0.510 | 0.601 |
| | Baseline (RoBERTa) | 0.640 | **0.558** | 0.643 |
| | Enhanced (RoBERTa) | 0.649 | 0.550 | **0.645** |
| Gen. LLM Prompting | Zero-Shot | 0.250 | 0.258 | 0.299 |
| | CoT | 0.227 | 0.190 | 0.243 |
| | CoT + Table | 0.588 | 0.367 | 0.542 |
| | CoT + Table + RAG | 0.578 | 0.357 | 0.529 |
| | CoT + Table + Tone | 0.667 | 0.384 | 0.609 |
| | One-Shot | 0.234 | 0.229 | 0.124 |
| | Three-Shot | 0.273 | 0.273 | 0.157 |
| | Nine-Shot | 0.266 | 0.262 | 0.159 |
| | Dynamic | 0.386 | 0.371 | 0.409 |
| | Nine-Shot + Tone | 0.344 | 0.361 | 0.286 |
| Fine-tuned LLM | Direct | **0.698** | 0.474 | 0.632 |

For encoder models, the Baseline configurations perform slightly better than their Enhanced counterparts in Macro-F1, especially in DistilBERT. This behaviour is attributable to the nature of the 3-class task: at this coarse level of granularity, the class distribution is relatively balanced, so Smart Resampling's majority-class downsampling probably removes useful training signal without providing a compensating benefit. Among prompting approaches, taxonomy injection (*CoT + Table*) is the decisive step, nearly doubling Macro-F1 over unguided CoT (0.367 vs. 0.190). Rhetorical tone conditioning (*CoT + Table + Tone*) then further pushes accuracy to 0.667. The fine-tuned LLM achieves the best accuracy overall (0.698), but its Macro-F1 (0.474) remains below the encoder baselines, reflecting conservative predictions biased toward the majority class *Ambivalent*.

### 4.2 Task 2 (Direct Evasion)

Table 3 reports performance on the 9-class evasion strategy task, the hardest setting due to severe class imbalance (e.g., *Explicit* at 30.5% vs. *Partial/half-answer* at 2.3%) and the high semantic overlap among rhetorical strategies.

In this setting, the trend observed in Task 1 is reversed, with Enhanced encoders outperforming Baseline models. Specifically, DistilBERT and RoBERTa achieve Macro-F1 improvements of 8.1% and 7.5%, respectively. The Enhanced pipeline provides a clear and consistent advantage across both architectures: T5-generated paraphrases populate underrepresented evasion classes

Table 3: Task 2 (Direct Evasion) results for all systems.

| Model Family | Configuration | Acc. | Macro-F1 | Weighted-F1 |
|---|---|---|---|---|
| Encoders | Baseline (DistilBERT) | 0.260 | 0.211 | 0.193 |
| | Enhanced (DistilBERT) | 0.237 | 0.292 | 0.217 |
| | Baseline (RoBERTa) | 0.308 | 0.233 | 0.258 |
| | Enhanced (RoBERTa) | 0.308 | **0.308** | 0.285 |
| Gen. LLM Prompting | Zero-Shot | 0.127 | 0.118 | 0.097 |
| | CoT | 0.150 | 0.169 | 0.123 |
| | CoT + Table | 0.166 | 0.155 | 0.161 |
| | CoT + Table + RAG | 0.189 | 0.141 | 0.162 |
| | CoT + Table + Tone | 0.146 | 0.116 | 0.144 |
| | One-Shot | 0.117 | 0.117 | 0.166 |
| | Three-Shot | 0.302 | 0.233 | 0.284 |
| | Nine-Shot | **0.386** | 0.288 | **0.371** |
| | Dynamic | 0.247 | 0.197 | 0.248 |
| | Nine-Shot + Tone | 0.370 | 0.269 | 0.354 |
| Fine-tuned LLM | Direct | 0.263 | 0.220 | 0.188 |

Table 4: Mapping from Evasion to Clarity results for all systems.

| Model Family | Configuration | Acc. | Macro-F1 |
|---|---|---|---|
| Encoders | Baseline (DistilBERT) | 0.429 | 0.420 |
| | Enhanced (DistilBERT) | 0.555 | 0.490 |
| | Baseline (RoBERTa) | 0.549 | 0.496 |
| | Enhanced (RoBERTa) | 0.640 | 0.568 |
| Gen. LLM Prompting | Zero-Shot | 0.679 | 0.407 |
| | CoT | 0.681 | 0.397 |
| | CoT + Table | 0.555 | 0.372 |
| | CoT + Table + RAG | 0.602 | 0.396 |
| | CoT + Table + Tone | 0.685 | 0.403 |
| | One-Shot | 0.234 | 0.285 |
| | Three-Shot | 0.695 | 0.545 |
| | Nine-Shot | 0.685 | **0.608** |
| | Dynamic | 0.558 | 0.460 |
| | Nine-Shot + Tone | 0.623 | 0.550 |
| Fine-tuned LLM | Direct | **0.721** | 0.571 |

with semantically diverse examples, while the injected tone feature supplies a strong discriminative prior that helps separate closely related strategies such as *Dodging* from *Declining to answer*.

Prompting approaches struggle severely with this fine-grained task; notably, extending the CoT prompt with the taxonomy table or rhetorical tone is counterproductive rather than helpful, suggesting that the additional context crowds out fine-grained in-context reasoning (best CoT Macro-F1: 0.169 with plain CoT).

In contrast with Task 1, the Dynamic approach yields a notable drop compared to the static Nine-Shot, attributable to cosine-based retrieval grouping examples by conversational topic rather than rhetorical strategy, thereby injecting conflicting evasion labels into the in-context context.

The fine-tuned LLM similarly underperforms at 0.220 Macro-F1.

### 4.3 Mapping from Evasion to Clarity

The mapping approach first predicts the nine evasion strategies and then converts them deterministically to the three clarity labels: *Explicit* → *Clear Reply*; *Declining to answer* and *Claims ignorance* → *Clear Non-Reply*; all remaining strategies → *Ambivalent*. Table 4 compares all systems under this protocol.

For encoder models, the mapping protocol consistently degrades performance relative to direct Task 1 training, a finding coherent with the observations of Thomas et al. (2024) on base-level encoder architectures. Classification errors in the intermediate nine-way step propagate to the final label, as shown by Baseline DistilBERT dropping from 0.545 to 0.420 Macro-F1 and Baseline RoBERTa from 0.558 to 0.496. However, Enhanced RoBERTa is an exception, achieving higher Macro-F1 under the mapping protocol (0.568) than direct training (0.550). This suggests that Smart Resampling and tone conditioning improve nine-class discriminability, enabling the model to capture fine-grained rhetorical patterns.

For prompting approaches, Zero-Shot already achieves a competitive 0.407 Macro-F1, and extending the CoT with a taxonomy table does not help, since it is strictly correlated to Task 2.

As illustrated in Figure 1, few-shot in-context learning exploits this alignment effectively, with the static Nine-Shot configuration reaching 0.608 Macro-F1, the highest result across all prompting strategies and tasks.

Overall, decoder-based approaches exhibit a clear improvement over the Direct Clarity setting when the evasion-based mapping strategy is applied, consistently with the conclusions of Thomas et al. (2024). The fine-tuned LLM achieves the best accuracy results under the Mapping approach. This gain arises because fine-grained evasion supervision guides the model toward more accurate coarse-level distinctions: after mapping, nine-class errors become correct macro-label predictions, indicating that the model captures useful rhetorical coarse structure even when it confuses fine-grained evasive subtypes.

## 5 Conclusion

In this paper, we explored computationally efficient methods for detecting political evasion, uncovering important trade-offs in how models process complex rhetorical strategies. Our most successful and reliable approach was the evasion-to-clarity mapping. By fine-tuning a 3B parameter model to identify fine-grained evasion techniques and mapping those to macro-level clarity labels, we provided the model with a much stronger learning signal than direct classification. Interestingly, this same mapping approach reduced the performance of our baseline encoders, as early classification errors simply carried over to the final prediction. We observed a similar trade-off with our data balancing techniques: while our Enhanced pipeline (using semantic resampling and tone injection) was essential for helping encoders learn the highly imbalanced 9-class task, it actually removed useful training data and lowered scores on the broader 3-class task. Finally, our exploratory methods highlighted the limitations of treating political discourse like standard text. Dynamic few-shot prompting struggled because it grouped examples by conversational topic rather than the actual evasive behavior. Likewise, decomposing multi-part questions through prompt chaining underperformed; evaluating each sub-question individually failed to capture the overall clarity of the complete answer, and errors multiplied across the pipeline stages. Future work should build on these insights by integrating rhetorical metadata directly into fine-tuning and designing retrieval systems that match text based on evasion strategy rather than semantic subject matter.

## References


AilsntuaLab. 2024. Qevasion: Political question evasion dataset. Version: commit-hash or date accessed; Accessed: 2026-01-05.

Peter Bull. 2003. *The Microanalysis of Political Communication: Claptrap and Ambiguity*. Routledge, London.

J. T. Dillon. 1990. *The Practice of Questioning*, 1 edition. Routledge.

Codruta Liliana Girlea. 2017. *Deception Detection in Dialogues*. Ph.D. thesis, University of Illinois at Urbana-Champaign.

Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, Anirudh Goyal, Anthony Hartshorn, Aobo Yang, Archi Mitra, Archie Sravankumar, Artem Korenev, Arthur Hinsvark, Aston Gregerson, and 2 others. 2024. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Sandra Harris. 1991. Evasive action: How politicians respond to questions in political interviews. In Paddy Scannell, editor, *Broadcast Talk*, pages 76–99. SAGE Publications, London.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. Lora: Low-rank adaptation of large language models. In *International Conference on Learning Representations*.

Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692*.

Colin Raffel, Noam Shazeer, Adam Roberts, Katherine Lee, Sharan Narang, Michael Matena, Yanqi Zhou, Wei Li, and Peter J. Liu. 2023. Exploring the limits of transfer learning with a unified text-to-text transformer. *Preprint*, arXiv:1910.10683.

Pranav Rajpurkar, Robin Jia, and Percy Liang. 2018. Know what you don't know: Unanswerable questions for squad. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 784–789, Melbourne, Australia. Association for Computational Linguistics.

Nils Reimers and Iryna Gurevych. 2019. Sentence-bert: Sentence embeddings using siamese bert-networks. *CoRR*, abs/1908.10084.

Konstantinos Thomas, Giorgos Filandrianos, Maria Lymperaiou, Chrysoula Zerva, and Giorgos Stamou. 2024. "I never said that": A dataset, taxonomy and baselines on response clarity classification. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 5204–5233, Miami, Florida, USA. Association for Computational Linguistics.

Jason Wei and Kai Zou. 2019. EDA: Easy data augmentation techniques for boosting performance on text classification tasks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 6382–6388, Hong Kong, China. Association for Computational Linguistics.


## A Exploratory Analysis of Prompt Chaining and Question Decomposition

In addition to our primary fine-tuning methodologies, we explored a generative prompt chaining pipeline for "multi-barrelled" queries, where a journalist asks multiple distinct questions in the same turn. We hypothesized that decomposing such questions into atomic components would improve grounding and, consequently, clarity detection accuracy.

### A.1 Methodology

The prompt chaining pipeline operated only at inference time and consisted of three stages:

1. **Question Splitting:** The original question was decomposed into atomic sub-questions using deterministic parsing, supplemented by an LLM-based splitter for ambiguous cases.

2. **Sub-question Evaluation:** Each sub-question was paired with the full political response. The fine-tuned LLM from Section 3.6 evaluated whether the response addressed each sub-question and assigned an evasion strategy.

3. **Aggregation:** Sub-question labels were aggregated into a unified prediction. A response was labeled *Explicit* only if all sub-questions were explicitly answered; mixed cases containing both explicit and evasive components were labeled *Partial/half-answer*; cases in which all components yielded a refusal were labeled *Clear Non-Reply*. These intermediate labels were then collapsed into the three Clarity classes used in the main experiments.

### A.2 Results and Error Analysis

Although decomposition is methodologically appealing, the prompt chaining approach degraded performance compared to direct evaluation. Table 5 reports the results.

| Model Paradigm | Acc. | Macro-F1 | Weighted-F1 |
|---|---|---|---|
| Direct Clarity | 0.698 | 0.474 | 0.632 |
| Prompt Chaining | 0.685 | 0.419 | 0.607 |

Table 5: Performance comparison of generative approaches on the 3-class Clarity task. The Chaining protocol exhibits a notable drop in Macro-F1 compared to direct clarity classification.

**Discussion of Performance Degradation.** The Prompt Chaining pipeline achieved a Macro-F1 score of 0.419, underperforming the standard Mapping approach (0.571) based on the original paper's method. Error analysis suggests three main causes:

- **Error Propagation Across Pipeline Stages:** Chaining increased the number of discrete decisions required at inference time: splitting the question, evaluating each sub-question, and aggregating the resulting labels. Errors or instabilities at any stage could therefore be amplified by the deterministic aggregation rules.

- **Mismatch with the Holistic Annotation Target:** The final Clarity label is assigned to the complete question–answer exchange, not to exhaustive coverage of every sub-question. By evaluating each component separately, the pipeline shifted the decision criterion from overall clarity to comprehensive sub-question coverage.

- **Rigid and Over-conservative Aggregation:** The aggregation rules treated sub-questions as equally important, although some components function as framing, elaboration, or rhetorical pressure rather than independent information requests. This penalized natural conversational responses, inflating *Ambivalent* predictions and reducing recall for *Clear Replies*.

Overall, question decomposition may be useful as an analytical aid, but our 3B-parameter models performed better when evaluating the holistic question–answer context directly through the Mapping approach, rather than relying on fragmented sub-question judgments and deterministic aggregation heuristics.